

Apache Kafka with Spring

1. Overview

In this tutorial, we'll cover Spring support for Kafka and the level of abstractions it provides over native Kafka Java client APIs. We will develop a Spring Boot app to demonstrate sending and receiving of messages in Kafka using spring-kafka.

Spring Kafka brings the simple and typical Spring template programming model with a *KafkaTemplate* and Message-driven POJOs via *@KafkaListener* annotation.

Lab Solution

Complete solution for this lab is available at github:

```
cd ~ && git clone https://github.com/fenago/kafka-spring-lab.git
```

```
cd kafka-spring-lab/spring-solution/spring-kafka
```

```
mvn clean compile
```

Open folder in IntelliJ after running above commands.

2. PreReq and Setup

Make sure that kafka and zookeeper is running:

We also need to add the *spring-kafka* dependency to our *pom.xml*:

```
<dependency>
  <groupId>org.springframework.kafka</groupId>
  <artifactId>spring-kafka</artifactId>
  <version>2.5.8.RELEASE</version>
</dependency>
```

Our example application will be a Spring Boot application.

3. Configuring Topics

As Kafka topics are not created automatically by default, this application requires that you create the following topics manually.

```
$ cd ~/kafka-training/kafka
```

```
$ bin/kafka-topics.sh --create --zookeeper localhost:2181 --replication-factor 1 --
partitions 1 --topic fenago
```

```
$ bin/kafka-topics.sh --create --zookeeper localhost:2181 --replication-factor 1 --
partitions 5 --topic partitioned
```

```
$ bin/kafka-topics.sh --create --zookeeper localhost:2181 --replication-factor 1 --
partitions 1 --topic filtered
```

```
$ bin/kafka-topics.sh --create --zookeeper localhost:2181 --replication-factor 1 --
partitions 1 --topic greeting
```

We need to add the *KafkaAdmin* Spring bean, which will automatically add topics for all beans of type *NewTopic*:

```
@Configuration
public class KafkaTopicConfig {

    @Value(value = "${kafka.bootstrapAddress}")
    private String bootstrapAddress;

    @Bean
    public KafkaAdmin kafkaAdmin() {
        Map<String, Object> configs = new HashMap<>();
        configs.put(AdminClientConfig.BOOTSTRAP_SERVERS_CONFIG, bootstrapAddress);
        return new KafkaAdmin(configs);
    }

    @Bean
    public NewTopic topic1() {
        return new NewTopic("baeldung", 1, (short) 1);
    }
}
```

4. Producing Messages

To create messages, we first need to configure a [ProducerFactory](#). This sets the strategy for creating Kafka [Producer](#) instances.

Then we need a [KafkaTemplate](#), which wraps a *Producer* instance and provides convenience methods for sending messages to Kafka topics.

Producer instances are thread safe. So, using a single instance throughout an application context will give higher performance. Consequently, *KafkaTemplate* instances are also thread safe, and use of one instance is recommended.

4.1. Producer Configuration

```
@Configuration
public class KafkaProducerConfig {

    @Bean
    public ProducerFactory<String, String> producerFactory() {
        Map<String, Object> configProps = new HashMap<>();
        configProps.put(
            ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
            bootstrapAddress);
        configProps.put(
            ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
            StringSerializer.class);
        configProps.put(
            ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
            StringSerializer.class);
        return new DefaultKafkaProducerFactory<>(configProps);
    }

    @Bean
```

```

    public KafkaTemplate<String, String> kafkaTemplate() {
        return new KafkaTemplate<>(producerFactory());
    }
}

```

4.2. Publishing Messages

We can send messages using the *KafkaTemplate* class:

```

@Autowired
private KafkaTemplate<String, String> kafkaTemplate;

public void sendMessage(String msg) {
    kafkaTemplate.send(topicName, msg);
}

```

The *send* API returns a *ListenableFuture* object. If we want to block the sending thread and get the result about the sent message, we can call the *getAPI* of the *ListenableFuture* object. The thread will wait for the result, but it will slow down the producer.

Kafka is a fast stream processing platform. Therefore, it's better to handle the results asynchronously so that the subsequent messages do not wait for the result of the previous message.

We can do this through a callback:

```

public void sendMessage(String message) {

    ListenableFuture<SendResult<String, String>> future =
        kafkaTemplate.send(topicName, message);

    future.addCallback(new ListenableFutureCallback<SendResult<String, String>>() {

        @Override
        public void onSuccess(SendResult<String, String> result) {
            System.out.println("Sent message=[" + message +
                "] with offset=[" + result.getRecordMetadata().offset() + "]);
        }

        @Override
        public void onFailure(Throwable ex) {
            System.out.println("Unable to send message=["
                + message + "] due to : " + ex.getMessage());
        }

    });
}

```

5. Consuming Messages

5.1. Consumer Configuration

For consuming messages, we need to configure a [ConsumerFactory](#) and a [KafkaListenerContainerFactory](#). Once these beans are available in the Spring bean factory, POJO-based consumers can be configured using [@KafkaListener](#) annotation.

[@EnableKafka](#) annotation is required on the configuration class to enable detection of [@KafkaListener](#) annotation on spring-managed beans:

```

@EnableKafka
@Configuration
public class KafkaConsumerConfig {

    @Bean
    public ConsumerFactory<String, String> consumerFactory() {
        Map<String, Object> props = new HashMap<>();
        props.put(
            ConsumerConfig.BootstrapServersConfig,
            bootstrapAddress);
        props.put(
            ConsumerConfig.GroupIdConfig,
            groupId);
        props.put(
            ConsumerConfig.KeyDeserializerClassConfig,
            StringDeserializer.class);
        props.put(
            ConsumerConfig.ValueDeserializerClassConfig,
            StringDeserializer.class);
        return new DefaultKafkaConsumerFactory<>(props);
    }

    @Bean
    public ConcurrentKafkaListenerContainerFactory<String, String>
        kafkaListenerContainerFactory() {

        ConcurrentKafkaListenerContainerFactory<String, String> factory =
            new ConcurrentKafkaListenerContainerFactory<>();
        factory.setConsumerFactory(consumerFactory());
        return factory;
    }
}

```

5.2. Consuming Messages

```

@KafkaListener(topics = "topicName", groupId = "foo")
public void listenGroupFoo(String message) {
    System.out.println("Received Message in group foo: " + message);
}

```

We can implement multiple listeners for a topic, each with a different group Id. Furthermore, one consumer can listen for messages from various topics:

```

@KafkaListener(topics = "topic1, topic2", groupId = "foo")

```

Spring also supports retrieval of one or more message headers using the [@Header](#) annotation in the listener:

```

@KafkaListener(topics = "topicName")
public void listenWithHeaders(
    @Payload String message,
    @Header(KafkaHeaders.RECEIVED_PARTITION_ID) int partition) {
    System.out.println(
        "Received Message: " + message
    );
}

```

```
        + "from partition: " + partition);
    }
}
```

5.3. Consuming Messages from a Specific Partition

Notice that we created the topic *baeldung* with only one partition.

For a topic with multiple partitions, however, a *@KafkaListener* can explicitly subscribe to a particular partition of a topic with an initial offset:

```
@KafkaListener(
    topicPartitions = @TopicPartition(topic = "topicName",
    partitionOffsets = {
        @PartitionOffset(partition = "0", initialOffset = "0"),
        @PartitionOffset(partition = "3", initialOffset = "0")}),
    containerFactory = "partitionsKafkaListenerContainerFactory")
public void listenToPartition(
    @Payload String message,
    @Header(KafkaHeaders.RECEIVED_PARTITION_ID) int partition) {
    System.out.println(
        "Received Message: " + message
        + "from partition: " + partition);
}
```

Since the *initialOffset* has been set to 0 in this listener, all the previously consumed messages from partitions 0 and 3 will be re-consumed every time this listener is initialized.

If we don't need to set the offset, we can use the *partitions* property of *@TopicPartition* annotation to set only the partitions without the offset:

```
@KafkaListener(topicPartitions
    = @TopicPartition(topic = "topicName", partitions = { "0", "1" })))
```

5.4. Adding Message Filter for Listeners

We can configure listeners to consume specific types of messages by adding a custom filter. This can be done by setting a [RecordFilterStrategy](#) to the *KafkaListenerContainerFactory*:

```
@Bean
public ConcurrentKafkaListenerContainerFactory<String, String>
    filterKafkaListenerContainerFactory() {

    ConcurrentKafkaListenerContainerFactory<String, String> factory =
        new ConcurrentKafkaListenerContainerFactory<>();
    factory.setConsumerFactory(consumerFactory());
    factory.setRecordFilterStrategy(
        record -> record.value().contains("World"));
    return factory;
}
```

We can then configure a listener to use this container factory:

```
@KafkaListener(
    topics = "topicName",
```

```

    containerFactory = "filterKafkaListenerContainerFactory")
    public void listenWithFilter(String message) {
        System.out.println("Received Message in filtered listener: " + message);
    }

```

In this listener, all the **messages matching the filter will be discarded**.

6. Custom Message Converters

So far, we have only covered sending and receiving Strings as messages. However, we can also send and receive custom Java objects. This requires configuring appropriate serializer in *ProducerFactory* and deserializer in *ConsumerFactory*.

Let's look at a simple bean class*,* which we will send as messages:

```

public class Greeting {

    private String msg;
    private String name;

    // standard getters, setters and constructor
}

```

6.1. Producing Custom Messages

In this example, we will use [JsonSerializer](#).

Let's look at the code for *ProducerFactory* and *KafkaTemplate*:

```

@Bean
public ProducerFactory<String, Greeting> greetingProducerFactory() {
    // ...
    configProps.put(
        ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
        JsonSerializer.class);
    return new DefaultKafkaProducerFactory<>(configProps);
}

@Bean
public KafkaTemplate<String, Greeting> greetingKafkaTemplate() {
    return new KafkaTemplate<>(greetingProducerFactory());
}

```

We can use this new *KafkaTemplate* to send the *Greeting* message:

```

kafkaTemplate.send(topicName, new Greeting("Hello", "World"));

```

6.2. Consuming Custom Messages

Similarly, let's modify the *ConsumerFactory* and *KafkaListenerContainerFactory* to deserialize the *Greeting* message correctly:

```

@Bean
public ConsumerFactory<String, Greeting> greetingConsumerFactory() {
    // ...
}

```

```

    return new DefaultKafkaConsumerFactory<>(
        props,
        new StringDeserializer(),
        new JsonDeserializer<>(Greeting.class));
}

@Bean
public ConcurrentKafkaListenerContainerFactory<String, Greeting>
greetingKafkaListenerContainerFactory() {

    ConcurrentKafkaListenerContainerFactory<String, Greeting> factory =
        new ConcurrentKafkaListenerContainerFactory<>();
    factory.setConsumerFactory(greetingConsumerFactory());
    return factory;
}

```

The spring-kafka JSON serializer and deserializer uses the Jackson library, which is also an optional Maven dependency for the spring-kafka project.

So, let's add it to our *pom.xml*:

```

<dependency>
  <groupId>com.fasterxml.jackson.core</groupId>
  <artifactId>jackson-databind</artifactId>
  <version>2.9.7</version>
</dependency>

```

Instead of using the latest version of Jackson, it's recommended to use the version that is added to the *pom.xml* of spring-kafka.

Finally, we need to write a listener to consume *Greeting* messages:

```

@KafkaListener(
    topics = "topicName",
    containerFactory = "greetingKafkaListenerContainerFactory")
public void greetingListener(Greeting greeting) {
    // process greeting message
}

```

When the application runs successfully, following output is logged on to console (along with spring logs):

Message received from the 'fenago' topic by the basic listeners with groups foo and bar

```

Received Message in group 'foo': Hello, World!
Received Message in group 'bar': Hello, World!

```

Message received from the 'fenago' topic, with the partition info

```

Received Message: Hello, World! from partition: 0

```

Message received from the 'partitioned' topic, only from specific partitions

```

Received Message: Hello To Partioned Topic! from partition: 0
Received Message: Hello To Partioned Topic! from partition: 3

```

Message received from the 'filtered' topic after filtering

Received Message in filtered listener: Hello Fenago!

Message (Serialized Java Object) received from the 'greeting' topic

Received greeting message: Greetings, World!!

7. Conclusion

In this lab, we covered Spring support for Apache Kafka. We took a brief look at the classes used for sending and receiving messages.

Before running the code, please make sure that Kafka server is running and that the topics are created manually.